

Bristol-Myers Squibb Molecular Translation 3rd place solution

Team: kyamaro
(KF + lyakaap + camaro)



About us


- Kazuki Fujikawa (KF)
 - Data scientist at DeNA
 - Analyse company's baseball team
 - Kaggle Grandmaster
- Shuhei Yokoo (Iyakaap)
 - Data scientist at DeNA
 - Analyse company's baseball team
 - Kaggle Grandmaster
- Daisuke Yamamoto(Camaro)
 - Kaggle Master

Solution: Overview

Phase1:

Image Captioning

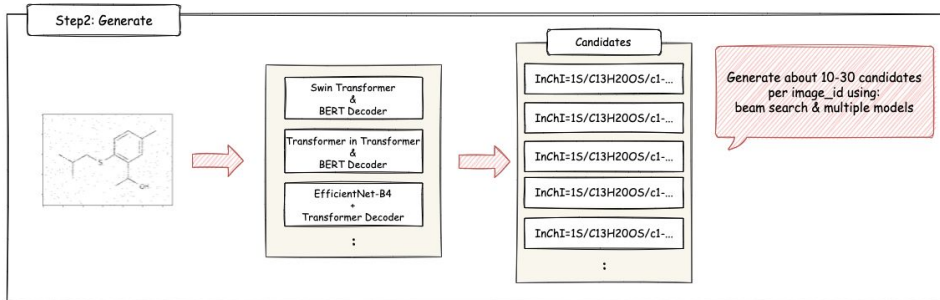
- Salt & pepper noise
- Focal loss



Phase2:

Generate InChI candidates

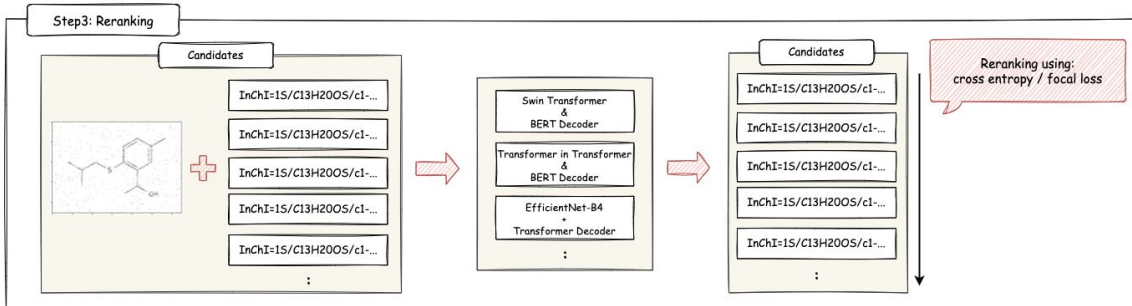
- Beam search
- Logit ensemble



Phase3:

Reranking InChI candidates

- Normalize with RDKit
- Rerank with multiple models




Solution: Overview

Phase1:

Image Captioning

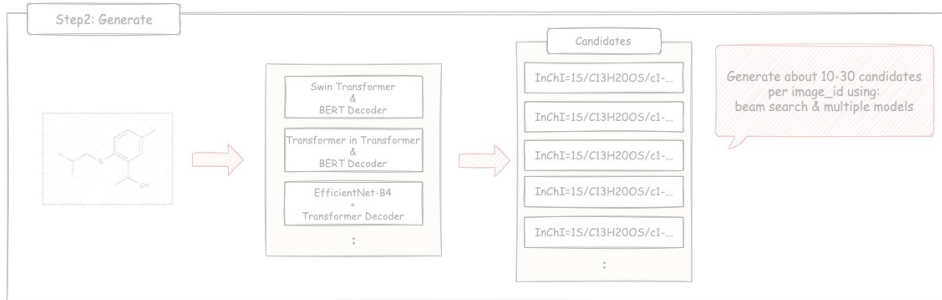
- Salt & pepper noise
- Focal loss



Phase2:

Generate InChI candidates

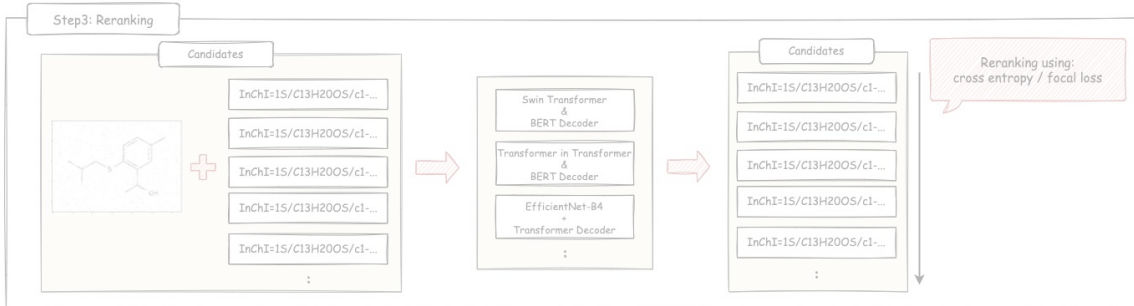
- Beam search
- Logit ensemble



Phase3:


Reranking InChI candidates

- Normalize with RDKit
- Rerank with multiple models



Phase1: Image captioning framework


- Generate sequences by considering InChI as a just string, in the framework of the general image captioning task
 - Surprisingly work well, although each layer has different meanings
 - Performance: Transformer based > CNN based (in our settings)
 - For this task, we need to capture "what" and "where" is in the image, so a Transformer-based Encoder would have been more advantageous.



Phase1: Denoise salt and pepper (LB: 2.18→1.26)

- In a certain size kernel, if there is no dot other than center, consider it as noise and remove it.
- Even without fine tuning, CV/LB gap was drastically improved.
- CV got slightly worse maybe because it sometimes removes an useful information too.


Example of denoise filter (k=5)



Model	CV	LB
TNT	1.13	2.18
+denoise (k=3)	1.69	1.71
+denoise (k=5)	1.22	1.26
+denoise (k=7)	1.17	1.32

Phase1: Add salt and pepper (LB: 1.26→1.06)


- CV/LB gap was also improved by adding salt and pepper noise at training time, even without denoise at test time.



Phase1: CE Loss → Focal Loss

- After training a few epochs, our model predicts surprisingly well and almost all samples got too easy one.
 - Ex. Swin Transformer (CV: 0.98, LB: 0.99)
 - Levenshtein=0: 87%
 - Levenshtein=1: 7%
 - ...
- Using Focal Loss makes training more efficient by focusing on hard example.

Hist of Levenshtein (in Log scale)




$$\text{FL}(p_t) = -(1 - p_t)^\gamma \log(p_t).$$

Solution: Overview

Phase1:

Image Captioning

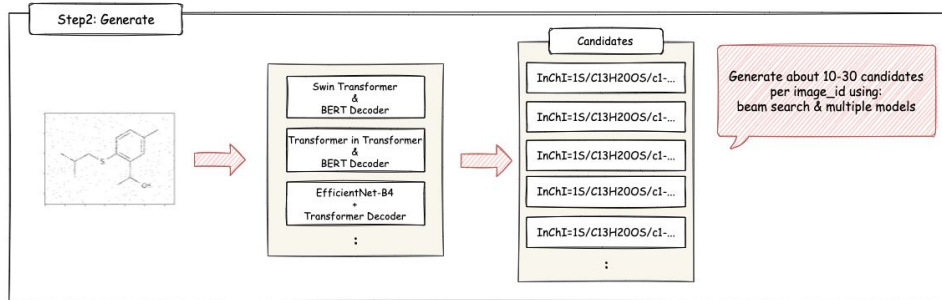
- Swin Transformer
- Salt & pepper noise
- Focal loss
- Train longer



Phase2:

Generate InChI candidates

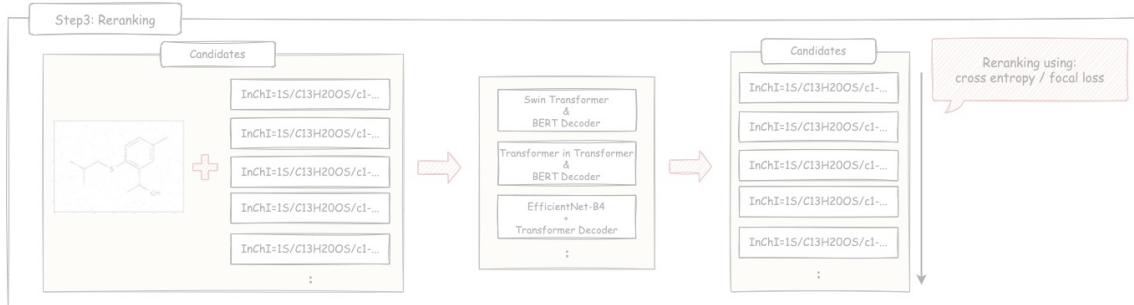
- Beam search
- Logit ensemble



Phase3:


Reranking InChI candidates

- Normalize with RDKit
- Rerank with multiple models



Phase2: Beam Search

- To make reranking at phase 3 more effective, need to produce diverse and high quality candidates.
 - Combining Beam Search and Greedy method was effective
 - Use broad Beam Search (ex.k=32) only for hard examples (Iyakaap)

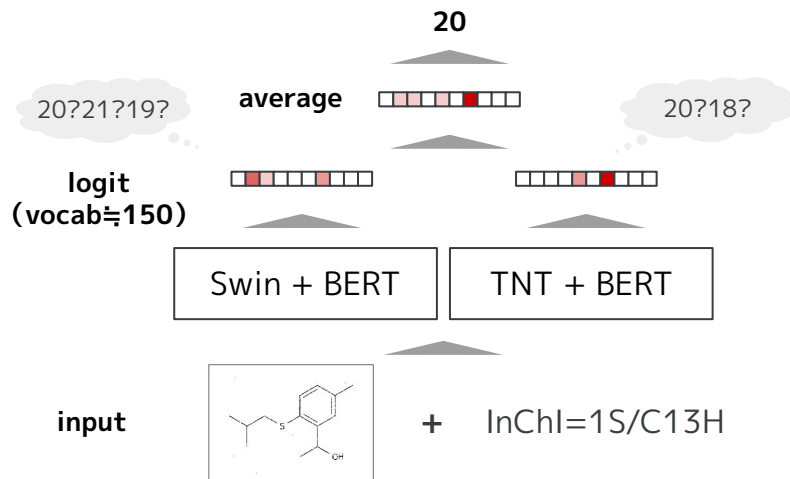


InChI=1S/C13H20OS...
InChI=1S/C13H21OS...
InChI=1S/C11H21OS...
InChI=1S/C12H20OS...

Model	CV
Swin (beam=1)	0.98
Swin (beam=4)	0.93
Swin (beam=1+4)	0.87

Phase2: Logit ensemble

- Ensemble logits at InChI generation phase
 - There is a difference in predictions even between Swin and TNT, so ensemble was effective.
 - But it was difficult to mix all of 3 members' models due to implementation difference.



Model	CV
Swin (beam=1)	0.98
TNT (beam=1)	1.04
Swin+TNT (beam=1)	0.83

Solution: Overview

Phase1:

Image Captioning

- Swin Transformer
- Salt & pepper noise
- Focal loss
- Train longer

Phase2:

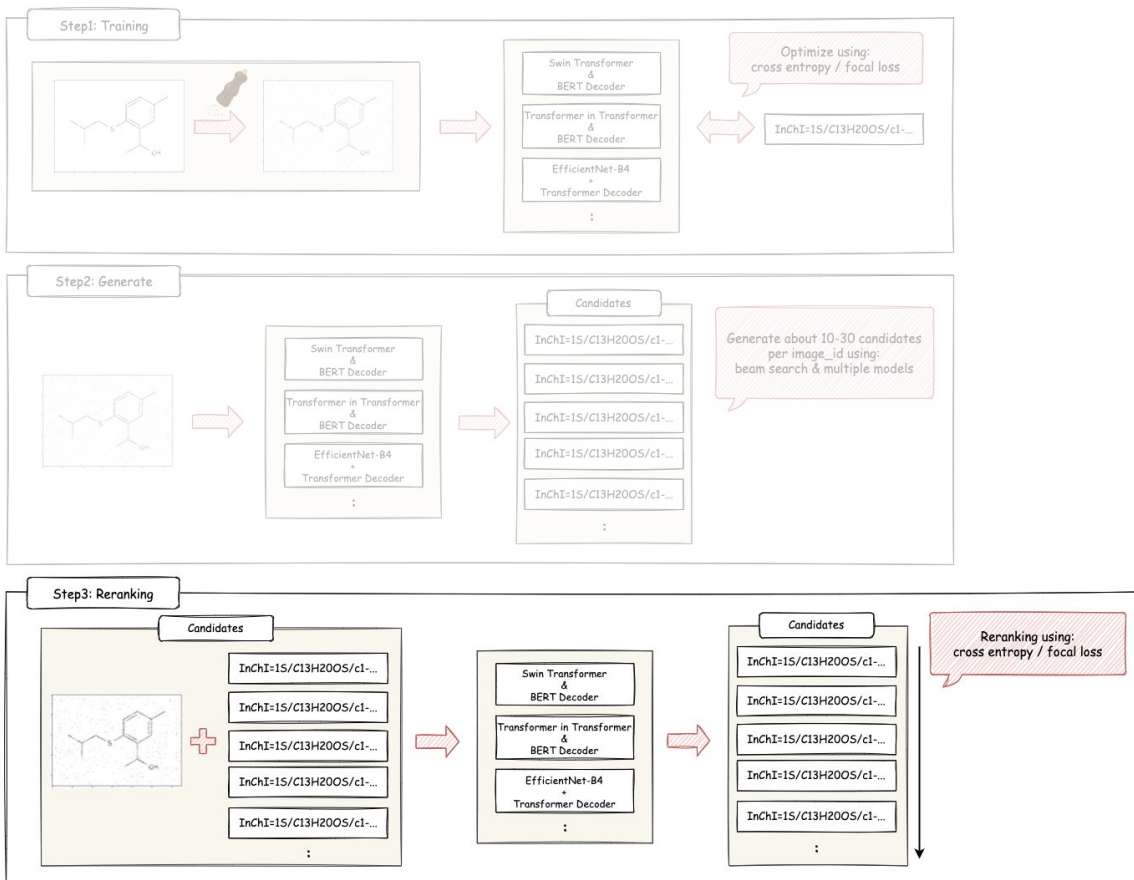
Generate InChI candidates

- Beam search
- Logit ensemble

Phase3:

Reranking InChI candidates

- Normalize with RDKit
- Rerank with multiple models




Phase3: RDKit normalization

- Rerank candidates with using RDKit normalization
 - Even though our model choose the most likely token in InChI generation, still there are many tiny mistakes.
 - RDKit normalization was effective way to fix such mistakes, but sometimes it makes score worse.
 - We re-score normalized candidates by calculating loss between generated candidates and teacher-forced model predictions.
 - ※ loss is either cross entropy loss or focal loss
 - Then rerank by below logic:

```
df = df.sort_values(  
    by=["is_valid", "score"],  
    ascending=[False, True],  
).groupby("image_id").first()
```

Phase3: Rerank by multiple models each other

- Gather candidates from models, then evaluate by themselves
 - Each model calculates loss and take average
 - Take rank average between the model groups which have a gap in obsolete value range
 - It was easy to apply to diverse models because it doesn't mind implementation difference



InChI Candidates	Model (CE)				Model (Focal)				Average
	A	B	Mean	Rank	C	D	Mean	Rank	Rank
InChI=1S/C13H20OS...	0.1	0.9	0.5	3	0.03	0.05	0.04	2	2.5
InChI=1S/C13H21OS...	0.2	0.2	0.2	1	0.02	0.02	0.02	1	1
InChI=1S/C12H20OS...	0.3	0.5	0.4	2	0.01	0.09	0.05	3	2.5

※ In this example we take rank average only in same image id, but actually we took globally.

Final models (part of)

Member	Model	CV	LB (greedy)	LB (reranked)
KF	Swin large (384x384) + PL	0.90	0.89	0.79
	TNT (512x1024) + PL	0.97	0.99	0.87
Iyakaap	Swin base (384x384) + PL	0.92	0.85	0.74
	Swin base (384x384) + PL + focal loss	0.87	0.78	0.70
	EffNet-v2 >> ViT (448x448) + PL + focal loss	0.86	0.81	0.70
Camaro	EffNetB4 Transformer Decoder (416x736)	0.67	0.87	0.67
	EffNetB4 Transformer Decoder (416x736) + PL	0.67	0.73	0.62
	EffNetB4 Transformer Decoder (416x736) + Noise	0.65	0.84	0.62
	EffNetB4 Transformer Decoder (416x736) + Noise/Denoise	0.83	0.77	0.61



3



kyamaro



0.53

90

4d

Summary

- It was the key to identify the cause of Train/Test Gap
 - It seems not so many top team found it?
- Generate many candidates by diverse models + rerank
 - The more candidates or the more models, score is the better
 - After team merging(kyakaap + camaro), score was drastically improved. Diversity wins.
- Beam search + Greedy was effective way to increase candidate
 - Beam search itself should contain greedy candidates, but using both was better somehow. Maybe due to too strong pruning?
 - Other smart search(like MCTS) may effective, though it was hard to apply due to hardware limit.

Discussions

- Why test data is a bit noisier than train data? (Is it intentional?)
- How to make use of invalid prediction in terms of InChI rules?
- How much was original expectation for winning score?
- How to merge solutions from 3 teams?
- What's next for this competition?
 - Actually digitize old publications? Write paper?

Appendix


Phase1: lyakaap part



- Encoder:
 - Swin-transformer (size=384x384)
 - EfficientNet-v2 + ViT (size=448x448)
 - CNN as Patch embedding
- Decoder:
 - 3 layers transformer
 - Increasing number of layers more makes training unstable
- PIL.Image.resize is better than cv2.resize or other resizer ([reference](#))
- Fine tuning with pseudo labeling
 - It improves score with reducing CV/LB gap
- Training speed up by Sequence bucketing

Baseline: RDKit normalization (postprocess)

- Normalize generated InChIs by RDkit ([by nofreewill](#))
 - Execute MolToInchi(MolFromInchi(inchi))
 - This process can correct inchis if the atoms are just mis-numbered



InChI=1S
/C9F5NO
/c10-4-5(11)7(13)9(15-2-1-3-16)8(14)6(4)12



InChI=1S
/C9F5NO
/c10-4-5(11)7(13)9(8(14)6(4)12)15-2-1-3-16

AGENDA


- Baseline
- **Solution**

Phase1: Compare predictions between Swin and TNT




- Swin has CV-LB better CV/LB relation ship than TNT
- To dig in the gap, calculate Levenshtein between these models' predictions. Even between models, there was a gap (val: 1.02 test:2.16)
- Check images which have large Levenshtein both in valid test
 - Tend to have large image size both in val/test
 - Only in test, tend to have much salt and pepper noise

Worst Levenshtein (valid)



Worst Levenshtein (test)



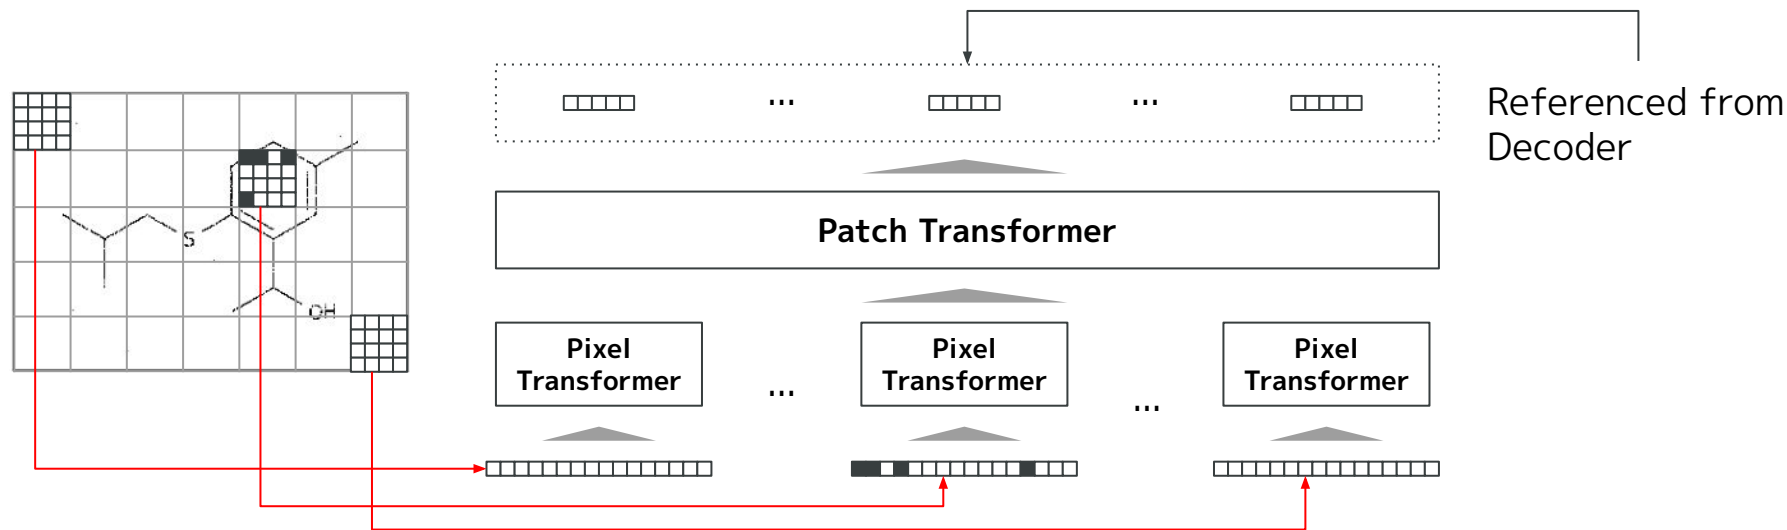
Phase1: Train longer



- There was a difference between kyakaap and camaro in CV/LB relationship.
 - Best single (Camaro): CV=0.66, LB=0.87
 - Best single (KF): CV=0.98, LB=0.99
- It can be due to train steps.
 - Camaro: 25-50 epoch (TPU 7days?)
 - KF: 10 epoch (A100 7days)
- After training kyakaap model 3 epochs more, CV improved from 0.98 to 0.91

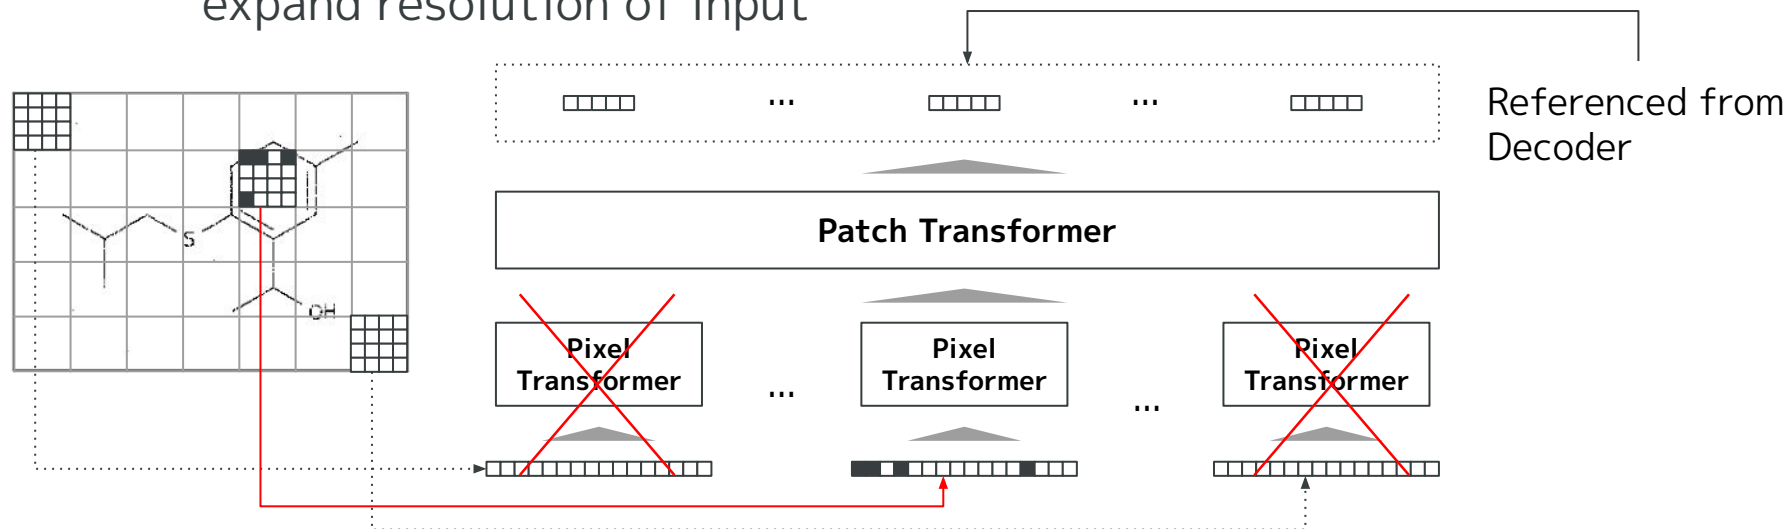
Phase1: TNT Encoder

- Use Transformer in Transformer (TNT) as Encoder
 - Reported this works well [by hengck23](#)
 - TNT encodes the relationships among pixels by Transformer, unlike ViT which encodes these by linear mapping



Phase1: TNT Encoder with variable patches (CV:1.30, LB:2.35)

- Majority of the areas are empty in images from this competition
 - Apply TNT on only patches where not empty. ([by hengck23](#))
 - → Reduced computational cost and GPU memory, making it easier to expand resolution of input



Phase1: Swin Transformer Encoder (LB:2.18→LB:1.43)




- Swin transformer has good CV-LB gap (CV: 1.23, LB: 1.43)
 - Better than TNT (CV: 1.30, LB: 2.18)
 - (We don't have a enough explanation for this though...)

提案手法: Swin Transformer

Window based Multihead Self-Attention (W-MSA)

- 画像をパッチに分割後、パッチの集合であるウィンドウを定義
- Window内のパッチに対してのみ、Self-Attentionで参照する
 - → Self-Attentionの計算コストは画像サイズの大きさに対して線形に増加




12

提案手法: Swin Transformer

Shifted window based Multihead Self-Attention (SW-MSA)

- W-MSA では、ウィンドウ間の関係性をモデリングできない
 - → ウィンドウをシフトさせ、ウィンドウ間関係性をモデリングできるようにした
 - (下図: 縦方向に2patch, 横方向に2patch, ウィンドウをシフトしている)



13

<https://www.slideshare.net/DeepLearningJP2016/dlswin-transformer-hierarchical-vision-transformer-using-shifted-windows>